

1. Nearest Neighbors:

1.6.1. Unsupervised Nearest Neighbors

`NearestNeighbors` implements unsupervised nearest neighbors learning. It acts as a uniform interface to three different nearest neighbors algorithms: `BallTree`, `KDTree`, and a brute-force algorithm based on routines in `sklearn.metrics.pairwise`. The choice of neighbors search algorithm is controlled through the keyword `'algorithm'`, which must be one of `['auto', 'ball_tree', 'kd_tree', 'brute']`. When the default value `'auto'` is passed, the algorithm attempts to determine the best approach from the training data. For a discussion of the strengths and weaknesses of each option, see [Nearest Neighbor Algorithms](#).

metric : string or callable, default 'minkowski'

metric to use for distance computation. Any metric from scikit-learn or `scipy.spatial.distance` can be used.

If metric is a callable function, it is called on each pair of instances (rows) and the resulting value recorded. The callable should take two arrays as input and return one value indicating the distance between them. This works for Scipy's metrics, but is less efficient than passing the metric name as a string.

```
from sklearn.neighbors import NearestNeighbors
```

```
X = np.array([[-1, -1], [-2, -1], [-3, -2], [1, 1], [2, 1], [3, 2]])
```

```
nbrs = NearestNeighbors(algorithm='brute').fit(X)
```

```
nbrs
```

```
NearestNeighbors(algorithm='brute', leaf_size=30, metric='minkowski',  
                 metric_params=None, n_jobs=None, n_neighbors=5, p=2, radius=1.0)
```

```
nbrs.kneighbors(X)
```

```
(array([[0.          , 1.          , 2.23606798, 2.82842712, 3.60555128],  
       [0.          , 1.          , 1.41421356, 3.60555128, 4.47213595],  
       [0.          , 1.41421356, 2.23606798, 5.          , 5.83095189],  
       [0.          , 1.          , 2.23606798, 2.82842712, 3.60555128],  
       [0.          , 1.          , 1.41421356, 3.60555128, 4.47213595],  
       [0.          , 1.41421356, 2.23606798, 5.          , 5.83095189]]),  
 array([[0, 1, 2, 3, 4],  
       [1, 0, 2, 3, 4],  
       [2, 1, 0, 3, 4],  
       [3, 4, 5, 0, 1],  
       [4, 3, 5, 0, 1],  
       [5, 4, 3, 0, 1]], dtype=int64))
```

2. Join two DataFrame:

```
df = pd.DataFrame({"key": ["K0", "K1", "K2", "K3", "K4", "K5"],
                  "A": ["A0", "A1", "A2", "A3", "A4", "A5"]})
```

df

	A	key
0	A0	K0
1	A1	K1
2	A2	K2
3	A3	K3
4	A4	K4
5	A5	K5

```
other = pd.DataFrame({"key": ["K0", "K1", "K2"],
                      "B": ["B0", "B1", "B2"]})
```

other

	B	key
0	B0	K0
1	B1	K1
2	B2	K2

```
df.join(other, lsuffix="_caller", rsuffix="_other")
```

	A	key_caller	B	key_other
0	A0	K0	B0	K0
1	A1	K1	B1	K1
2	A2	K2	B2	K2
3	A3	K3	NaN	NaN
4	A4	K4	NaN	NaN
5	A5	K5	NaN	NaN

```
df.set_index("key").join(other.set_index("key"))
```

	A	B
key		
K0	A0	B0
K1	A1	B1
K2	A2	B2
K3	A3	NaN
K4	A4	NaN
K5	A5	NaN

```
df.join(other.set_index("key"), on="key")
```

	A	key	B
0	A0	K0	B0
1	A1	K1	B1
2	A2	K2	B2
3	A3	K3	NaN
4	A4	K4	NaN
5	A5	K5	NaN

3. Rename a column:

```
jack_words.rename(columns = {"count": "count - jack"})
```

	word	count - jack
141	the	24
7	and	15
101	of	12
142	theatre	7
62	for	6

4. Convert a collection of raw documents to a matrix of TF-IDF features:

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
corpus = [
    "This is the first document.",
    "This document is the second document.",
    "And this is the third one.",
    "Is this the first document?"
]
```

```
vectorizer = TfidfVectorizer()
```

```
X = vectorizer.fit_transform(corpus)
```

```
X
```

```
<4x9 sparse matrix of type '<type 'numpy.float64'>'
  with 21 stored elements in Compressed Sparse Row format>
```

```
X.shape
```

```
(4, 9)
```

```
print vectorizer.get_feature_names()
```

```
[u'and', u'document', u'first', u'is', u'one', u'second', u'the', u'third', u'this']
```

```
X.toarray()

array([[0.          , 0.46979139, 0.58028582, 0.38408524, 0.          ,
        0.          , 0.38408524, 0.          , 0.38408524],
       [0.          , 0.6876236 , 0.          , 0.28108867, 0.          ,
        0.53864762, 0.28108867, 0.          , 0.28108867],
       [0.51184851, 0.          , 0.          , 0.26710379, 0.51184851,
        0.          , 0.26710379, 0.51184851, 0.26710379],
       [0.          , 0.46979139, 0.58028582, 0.38408524, 0.          ,
        0.          , 0.38408524, 0.          , 0.38408524]])
```

5. Generate a random array with 5 rows and 3 columns:

```
np.random.seed(0)

np.random.randn(5, 3)

array([[ 1.76405235,  0.40015721,  0.97873798],
       [ 2.2408932 ,  1.86755799, -0.97727788],
       [ 0.95008842, -0.15135721, -0.10321885],
       [ 0.4105985 ,  0.14404357,  1.45427351],
       [ 0.76103773,  0.12167502,  0.44386323]])
```

6. Convert True/False to 1/0:

```
doc.dot(random_vectors) >= 0

array([[False,  True, False, False,  True,  True, False,  True, False,
        True, False, False, False,  True,  True,  True]])

np.array(doc.dot(random_vectors)>=0, dtype=int)

array([[0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1]])
```

7. `itertools.combinations(iterable, r)`: Return r length subsequences of elements from the input iterable.

```
from itertools import combinations

num_vector = 16
search_radius = 3

for diff in combinations(range(num_vector), search_radius):
    print diff

(0, 1, 2)
(0, 1, 3)
(0, 1, 4)
(0, 1, 5)
(0, 1, 6)
(0, 1, 7)
(0, 1, 8)
(0, 1, 9)
(0, 1, 10)
(0, 1, 11)
(0, 1, 12)
(0, 1, 13)
(0, 1, 14)
(0, 1, 15)
```